



# PROGRAMMING ASSIGNMENT # 2 – SUPPLEMENTARY DISCUSSION

Tanzir Ahmed  
CSCE 313 Fall 2018



# Programming Assignment # 2

- Main Challenges (Still after fork + exec)
  - *I/O Redirection*
  - *Piping*
- To solve these challenges, we need the topics:
  - *File I/O basics*
  - *Inter Process Communication using **pipes***

# File I/O – File Descriptors

```
int main ()
{
    int fd; // file descriptor
    char* buf [] = "file content";
    fd = open ("foobar.txt", O_CREAT|O_WRONLY);
    write (fd, buf, strlen (buf)); close (fd);
    fd = open("foobar.txt", O_RDONLY, 0);
    read(fd, &c, 1);
    printf("c=%c\n", c);
    close (fd);
    return 0;
}
```

Every process has a file descriptor table, where there are 3 default entries to begin with:

- *Standard input, output, and error*

**Descriptor table**

[one table per process]

stdin	fd 0	
stdout	fd 1	
stderr	fd 2	
	fd 3	fd

# I/O Redirection

Question: How does a shell implement I/O redirection?

```
unix> ls > foo.txt
```

Answer: By calling the **dup2(source, destination)** function

- Copies (per-process) descriptor table entry **source** to entry **destination**

Descriptor table

*before* dup2 (4, 1)

fd 0	
fd 1	a
fd 2	
fd 3	
fd 4	b



Descriptor table

*after* dup2 (4, 1)

fd 0	
fd 1	b
fd 2	
fd 3	
fd 4	b

# Implementing “ls -la>foo.txt”

```
#include <stdio.h>
#include <unistd.h>
#include <sys/stat.h>
#include <fcntl.h>
int main ()
{
    int fd = open ("foo.txt", O_CREAT|O_WRONLY,
        S_IRUSR | S_IWUSR);
    dup2 (fd, 1); // overwriting stdout with the new file
    execlp ("ls", "ls", "-l", "-a", NULL); // now execute
    return 0;
}
```


# IPC Pipe

BEFORE pipe



Process has some usual files open

AFTER pipe



Kernel creates a pipe and sets file descriptors

# IPC Pipe - Method

```
#include <stdio.h>
#include <unistd.h>


void main ()
{
    char buf [10];
    int fds [2];
    pipe (fds);
    printf ("sending msg: Hi\n");
    write (fds[1], "Hi", 3);
    read (fds[0], buf, 3);
    printf ("Received msg: %s\n", buf);
}
```

Connects the  
two fds as pipe

```
compute-linux1 tanzir/code> ./a.out
sending msg: Hi
Received msg: Hi
```

# Pipe Between Two Processes

```
int main ()
{
    int fds [2];
    pipe (fds); // connect the pipe
    if (!fork()){ // on the child side
        char * msg = "a test message";
        printf ("CHILD: Sending %s\n",
msg);
        write (fds [1], msg,
strlen(msg)+1);
    }else{
        char buf [100];
        read (fds [0], buf, 100);
        printf ("PARENT: Received %s\n",
buf);
    }
    return 0;
}
```





# Shell Piping Example:

**“ls -l | grep soda”**

■ Meaning of the command:

- *Find all files that has the string “soda” in the filename and show detailed properties of those files*

■ How many processes do we have to run (in addition to our shell process)?

- *Process # 1: To run “ls -l”*
- *Process # 2: To run “grep soda”*

■ What else do we need so that the process #1 sends its output to process #2

- *Idea: If we can connect **stdout of p1** to **stdin of p2**, we are done!!*
- *Step 1: Redirect stdout of p1 to a file descriptor fd1*
- *Step 2: Redirect stdin of p2 to a another file descriptor fd2*
- *Step 3: Now, pipe fd1 and fd2 together so that fd1 is the “write side” and fd2 is the “read side”*

# Shell Piping:

`“ls -l | grep soda”`

```
void main ()
{
    int fds [2];
    pipe (fds); // connect the pipe
    if (!fork()) { // on the child side
        dup2 (fds[1], 1); // redirect stdout to pipeout
        execlp ("ls", "ls", "-l", NULL);
    } else {
        dup2 (fds[0], 0); // redirect stdin to pipe in
        execlp ("grep", "soda", "-l", NULL);
    }
}
```

Step 3

Step 1

Step 2

# Difficulty with the Previous

- You have to hard code the stages
  - *2 pipes (3 pipe separated portions) will change the code completely: cannot generalize this code*
  - *Will require us to have multiple pairs of file descriptors*
- Can we think of a more general way of doing the same?

`ls -l | grep this | grep that | ...`

Child process  
(redirects  
stdout)

Parent process  
(redirects  
stdin)

# A General Pipe Portion

```
int fd [2];
pipe (fd);
if (fork() == 0){
    dup2 (fd[1], 1); // overwriting stdout to the pipe's WRITE end
    close (fd[0]); // close unused pipe end
    execute (portion[i]);
}else{
    wait (0);
    close (fd[1]); // close unused pipe end
    dup2 (fd [0],0); // overwriting stdin to pipe's READ end for the later portions
}
```

the later portions:

- Shouldn't the #of fd's be  $2 * (\text{\#of portions})$
- Also, should the last portion do the same thing (i.e., redir its stdout)?
- You also need to close the unused pipe ends
  - *Otherwise, your program will deadlock*

# Over All – Repeat the Following

- Read the line full of command
- Split the line into portions by “|”
- For each portion except the last:
  - *Pipe()*
  - *Fork() a child*
  - *Redirect the child's STDOUT to Write End of pipe*
  - *Execute the current portion under this child*
  - *Under the parent, just redirect the STDIN to the Read End of the pipe*